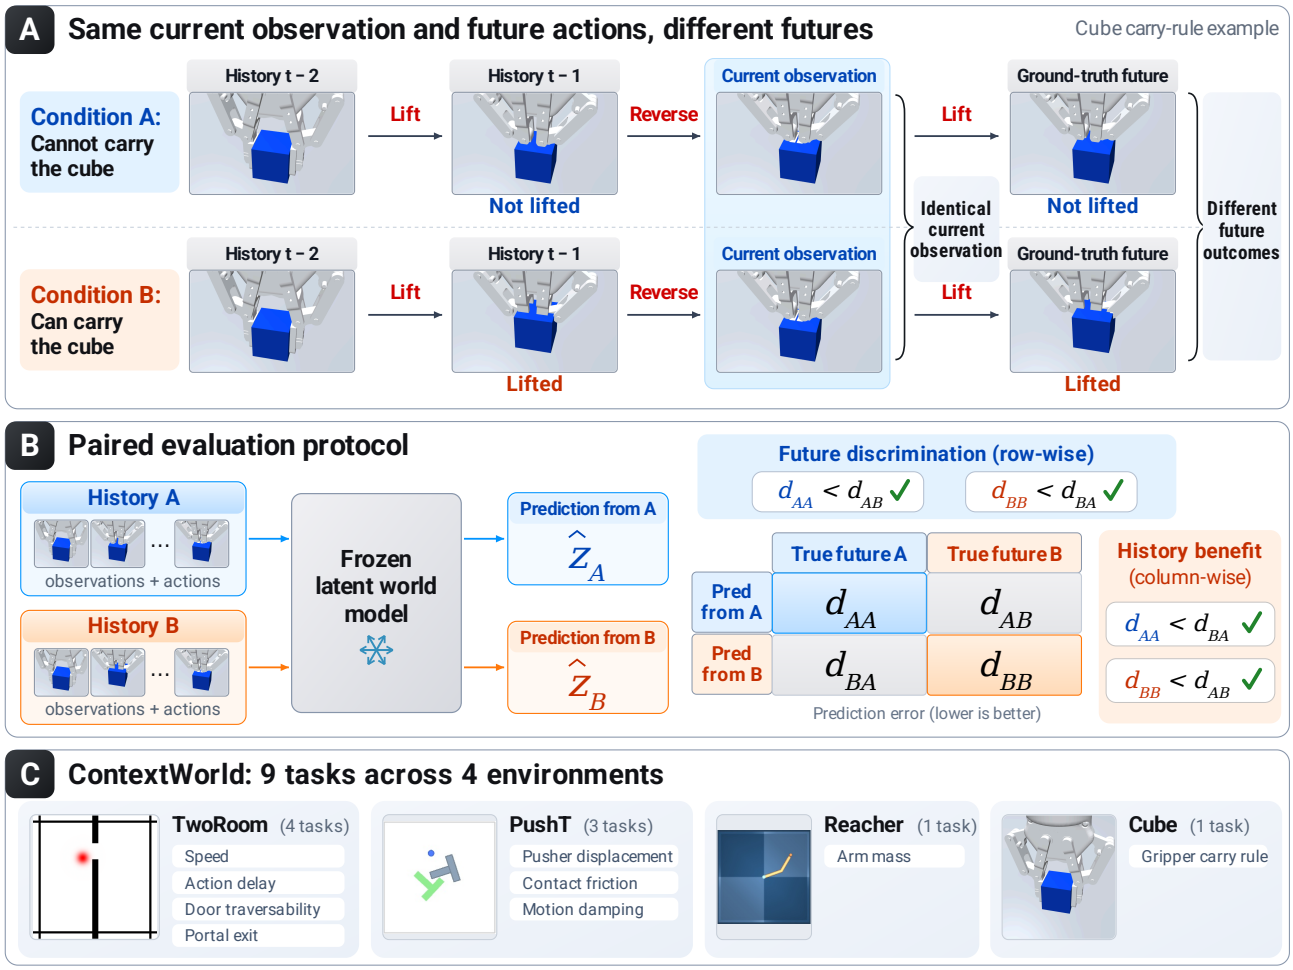


Figure 1. ContextWorld overview. (A) Recorded Cube carry-rule states, re-rendered at high resolution. A lift command and its reverse reveal whether the cube is carried, then return both conditions to the same current observation. Repeating the lift produces different futures. Commands label transitions. (B) Paired evaluation with frozen weights. Given each history and the shared future actions, the model predicts a latent future. Rows of the error matrix compare candidate futures; columns compare histories for a fixed target. The inequalities illustrate correct discrimination and a benefit from matching history. (C) Nine tasks across four environments.

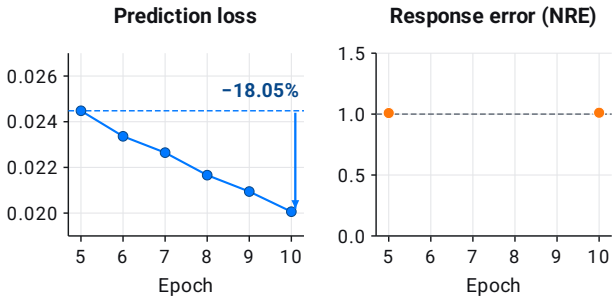


Figure 4. Prediction loss and conditional response. PLDM on contact friction, Original → Mixed (Joint). Recorded prediction loss falls by 18.05% from epochs 5 to 10, while training-query NRE stays near the zero-response reference of 1. Loss uses mixed training batches; NRE is measured on fixed paired training queries at the two endpoints. Full task and scheme comparisons remain in the manuscript appendix.

Prediction Loss and Adaptation

Weak evaluation performance may reflect poor generalization or conditional responses that were not learned well during training. We compare prediction loss with NRE on fixed training queries at epochs 5 and 10.

On friction, PLDM joint training reduces prediction loss by 18.05%, yet training-query NRE remains near the zero-response reference. On damping with the encoder frozen, loss approaches convergence while NRE improves but remains high. Changing the scheme helps: freezing yields a stronger conditional response on friction, whereas joint training performs better on damping.

These results expose a disconnect between prediction-loss optimization and history-conditioned adaptation, even on training data. Prediction loss aggregates errors across mixed data, whereas NRE directly measures error in the predicted differences between conditions.